

# English-to-English Grapheme-to-Phoneme Transliteration using RNN-Transducer

EE698R Course Project

Kumar Kanishk Singh (210544)  
Sunny Raja Prasad (218171078)  
Tanmey Agarwal (211098)

IIT Kanpur

April 20, 2025

# Outline

- 1 Introduction
- 2 Model Architecture
- 3 Evaluation
- 4 Conclusion

# Problem Statement I

- **Task:** English-to-English Grapheme-to-Phoneme (G2P) Transliteration
  - Converting written text (graphemes) to pronunciation (phonemes)
  - Critical component in text-to-speech and speech recognition systems
  - Challenging due to English's irregular spelling-to-sound correspondence
- **Input:** Sequence of English characters (graphemes)
  - Words like "knight", "phone", "through"
  - Each character processed sequentially through the Encoder
  - Represented as numerical IDs via character vocabulary mapping
- **Output:** Sequence of phonemes (pronunciations)
  - ARPABET phonetic notation (e.g., "N AY T" for "knight")
  - Variable-length sequences predicted by the model
  - Non-aligned with input (single grapheme may map to multiple phonemes)
- **Dataset:** CMU Pronouncing Dictionary (cmudict)
  - Contains 134,000+ English words and their pronunciations

# Problem Statement II

- Standard benchmark dataset for G2P tasks
- Available through NLTK library

- **Example:**

"knight" →

[N AY T]

"phone" →


[F OW N]

"through" →

[TH R UW]

# RNN-Transducer Architecture

- **Encoder:** Processes character sequence
- **Predictor:** Models previous phoneme predictions
- **Joiner:** Combines outputs of Encoder and Predictor
- **Loss:** RNN-Transducer loss (via torchaudio)



# Model Components

## **Encoder:** Character sequence processor

- Input: Tokenized character sequence  $(x_1, x_2, \dots, x_T)$
- Embedding layer:  $\text{dim} = 64$ , maps character IDs to vectors
- LSTM:  $\text{hidden\_dim} = 128$ , processes embeddings sequentially
- Handles variable-length inputs via packed sequences
- Output:  $f_t$  represents encoded character features

## **Predictor:** Phoneme context modeler

- Input: Previous phoneme sequence with blank token
- Embedding layer:  $\text{dim} = 64$ , maps phoneme IDs to vectors
- LSTM:  $\text{hidden\_dim} = 128$ , captures phoneme dependencies
- Output:  $g_u$  represents phoneme context features

# Model Components (Continued)

**Joiner:** Feature combination network

- Takes encoder output  $f_t$  and predictor output  $g_u$
- Projects both to  $\text{joint\_dim} = 256$  using linear layers:

$$h_{t,u} = \text{ReLU}(W_e f_t + W_p g_u)$$
$$\text{logits}_{t,u} = W_o h_{t,u}$$

- Creates a distribution over phoneme vocabulary
- Enables alignment between characters and phonemes

**Beam Search Decoder:** For efficient inference

- Maintains  $\text{beam\_width} = 3$  most probable hypotheses
- Iteratively evaluates encoder outputs at each timestep
- Handles the special blank symbol for RNN-T
- Algorithm:
  - 1 Initialize with empty hypothesis
  - 2 For each timestep  $t = 1 \dots T$ :
    - Extend each hypothesis with top-k predictions
    - Include blank transitions (stay on same alignment)
    - Keep top  $\text{beam\_width}$  hypotheses by score
  - 3 Return highest scoring complete hypothesis

## Data Source and Extraction

- CMU Pronouncing Dictionary via NLTK
- 134,000+ word-pronunciation pairs
- Example: ("knight", ["N", "AY", "T"])
- Convert all words to lowercase

## Vocabulary Construction

- Character vocabulary:  $|V_{\text{char}}| = 29$  (a-z, apostrophe, etc.)
- Phoneme vocabulary:  $|V_{\text{ph}}| = 40$  (ARPABET symbols)
- Special tokens:
  - <pad>| at index 0
  - <blank>| for RNN-T



## Numerical Conversion

- Convert text to indices:

"cat"  $\rightarrow$  [3, 1, 20]      [K, AE, T]  $\rightarrow$  [21, 3, 38]

## Batching and Padding

- Use custom `collate_fn` for batching
- Handle variable-length sequences with `pad_sequence`

## Dataset Setup

- 80% training set (107,200 samples)
- 20% test set (26,800 samples)
- Use `DataLoader` with batch size 64

## Optimization Parameters

- **Optimizer:** Adam
  - Default:  $\beta_1 = 0.9$ ,  $\beta_2 = 0.999$
  - No weight decay regularization
- **Learning Rate:**  $10^{-3}$  (fixed)
- **Batch Size:** 64
- **Training Epochs:** 5
- **Device:** CUDA if available, CPU as fallback

# Training Configuration: Loss and Loop

## Loss Function

- **RNN-T Loss** from `torchaudio.functional`:

$$\mathcal{L} = -\log P(y|x)$$

- Aligns predicted and target sequences via all valid paths
- Key steps:
  - 1 Apply `log_softmax` to logits
  - 2 Run forward-backward algorithm
  - 3 Handle blank tokens appropriately
  - 4 Compute batch mean loss

## Training Loop

- CER evaluation on training test sets
- Batch-level progress tracking
- Beam search decoding (width = 3)

# Evaluation Metric: Character Error Rate (CER)

## Definition:

$$\text{CER} = \frac{\text{EditDistance}(\text{ref}, \text{hyp})}{\text{Length of reference}}$$

## What is Edit Distance?

- Levenshtein distance between reference and hypothesis
- Measures:
  - **Substitution:** incorrect phoneme predicted
  - **Insertion:** extra phoneme added
  - **Deletion:** missing phoneme in prediction
- Normalized to [0–100%] range
- Lower CER  $\Rightarrow$  Better pronunciation prediction

# CER Example

**Reference:** [K AE T]

**Hypothesis:** [K AH T]

**Edit Distance:** 1 (substitution)

$$\text{CER} = \frac{1}{3} = 33.3\%$$

# Evaluation Process

- **Evaluation protocol:**

- ① Set model to evaluation mode (disable dropout)
- ② Process each test sample with beam search (width=3)
- ③ Convert predicted phoneme IDs to phoneme symbols
- ④ Compute CER between reference and hypothesis
- ⑤ Average CER across entire dataset

- **Reported metrics:**

- Training CER: 13.6% (model fit quality)
- Test CER: 14.4% (generalization ability)
- Transformer model comparison: 3.4% CER (both train/test)

# Error Analysis

## Common Error Patterns

- **Vowel substitutions** (most frequent)
  - Example: "about" → [AH B AW T] instead of [AH B AW T]
  - Cause: Similar phonetic properties between vowels
- **Silent letter handling**
  - Example: "knight" → [N AY T K] instead of [N AY T]
  - Cause: Model struggles with irregular spelling patterns
- **Complex phoneme mappings**
  - Example: "through" → [TH R U] instead of [TH R UW]
  - Cause: Multi-character graphemes to single phoneme mappings

## Error Distribution

- Substitutions: 68% of errors
- Deletions: 19% of errors
- Insertions: 13% of errors

# Results Comparison

Model	Train CER	Test CER
RNN-Transducer	13.6%	14.4%
Transformer	3.4%	3.4%

- Transformer models achieve significantly lower error rates (3.4% vs 14.4%)
- Better performance attributed to:
  - Deeper network architecture
  - Self-attention mechanism capturing global dependencies
  - Parallel processing of sequence elements
- RNN-T could potentially achieve similar results with:
  - Deeper LSTM layers
  - Larger hidden dimensions
  - But would require substantially more training time and compute resources




# Conclusion

- **Effective G2P Modeling:** The RNN-Transducer architecture demonstrates strong capability in learning grapheme-to-phoneme (G2P) mappings, effectively capturing dependencies between characters and phonemes.
- **Joint Encoder-Predictor Modeling:** By jointly modeling the character encoder and phoneme predictor, the RNN-T allows for dynamic alignment and better phoneme predictions, leading to improved performance over traditional sequence-to-sequence approaches.
- **Simplified Training with TorchAudio:** Leveraging the built-in RNN-T loss function from TorchAudio simplifies implementation and stabilizes training, reducing the need for complex loss function engineering.

- **Larger Beam Widths:** Investigate the impact of increasing beam width during decoding to explore more hypotheses and potentially improve final phoneme sequence quality.
- **Advanced LSTM Architectures:** Experiment with deeper LSTM stacks or bidirectional LSTMs in the encoder and predictor to better capture long-range and contextual dependencies in character and phoneme sequences.
- **Multilingual G2P Transliteration:** Extend the current RNN-T framework to handle multilingual datasets, enabling phoneme prediction across different languages and scripts for broader applicability in real-world transliteration tasks.

# References

-  M. Bisani and H. Ney.  
Joint-sequence models for grapheme-to-phoneme conversion.  
*Speech Communication*, 50(5):434–451, 2008.  
<https://dl.acm.org/doi/book/10.5555/2049202>
-  A. Graves.  
Sequence transduction with recurrent neural networks.  
In *ICML Representation Learning Workshop*, 2012.  
[https://www.cs.toronto.edu/~graves/icml\\_2012.pdf](https://www.cs.toronto.edu/~graves/icml_2012.pdf)
-  R. Fukuda, H. Yamamoto, and K. Takeda.  
Global RNN Transducer Models For Multi-dialect Speech Recognition.  
In *Interspeech*, 2022.  
[https://www.isca-archive.org/interspeech\\_2022/fukuda22\\_interspeech.pdf](https://www.isca-archive.org/interspeech_2022/fukuda22_interspeech.pdf)
-  Carnegie Mellon University.  
CMU Pronouncing Dictionary.  
[https://en.wikipedia.org/wiki/CMU\\_Pronouncing\\_Dictionary](https://en.wikipedia.org/wiki/CMU_Pronouncing_Dictionary)
-  Natural Language Toolkit.  
`nlk.corpus.reader.cmudict`.  
[https://www.nltk.org/\\_modules/nltk/corpus/reader/cmudict.html](https://www.nltk.org/_modules/nltk/corpus/reader/cmudict.html)
-  J. Jiang, Y. Zhang, S. Wang, et al.  
Pretraining Enhanced RNN Transducer.  
*Artificial Intelligence Research*, 2024.  
<https://www.sciopen.com/article/10.26599/AIR.2024.9150039>
-  H. Kang and K. Choi.  
An Ensemble of Grapheme and Phoneme for Machine Transliteration.  
*Proc. of the International Conference on Natural Language Processing*, 2005.  
<https://aclanthology.org/I05-1040.pdf>

Thank You!